AI/ML Applications in Biotechnology Research

Mini Project Proposal — Biotechnology × Artificial Intelligence

Sunny Lakhwani Biotechnology · Task 4: Mini-Project Proposal · 2026

Project Title	BioMind: An Integrated AI/ML Framework for Accelerated Drug Discovery and Genomic Insights
Domain	Medical Biotechnology × Artificial Intelligence / Machine Learning
Duration	12 months (Phase I: 6 months, Phase II: 6 months)
Team Size	4–6 members (Bioinformaticians, ML Engineers, Wet Lab Scientists)
Proposed Budget	USD $85,000 – $120,000 (research grant / university funding)

1. Introduction

The convergence of artificial intelligence (AI) and biotechnology is one of the most consequential scientific developments of our era. Biological datasets — genomic sequences, proteomics profiles, drug-target interaction networks, clinical trial outcomes — are growing at a rate that far outpaces human analytical capacity. Machine learning (ML) algorithms can identify patterns in these high-dimensional datasets with a speed and accuracy that traditional statistical methods cannot match.

This project proposes the development of **BioMind**, an integrated AI/ML framework designed to accelerate three critical bottlenecks in biotechnology: (1) novel drug target identification from multi-omics data, (2) protein structure prediction and ligand binding affinity optimization, and (3) clinical trial outcome prediction through patient stratification models. BioMind is designed as an open, modular platform accessible to academic and non-profit research institutions globally.

2. Problem Statement

Drug discovery and development remains one of the most expensive and time-consuming processes in science: it takes an average of 10–15 years and costs over $2.6 billion to bring a single drug to market, with a failure rate exceeding 90% in clinical trials (DiMasi et al., 2016).

Biotechnology research faces three compounding challenges:

- **Data Fragmentation:** Genomic, transcriptomic, proteomic, and clinical data are siloed across institutions with no unified analysis platform, leading to duplicated effort and missed biological insights.
- **Slow Target Identification:** Identifying causal disease drivers from thousands of differentially expressed genes requires manual expert curation — a bottleneck that takes months per disease indication.
- **High Clinical Attrition:** Most drugs that show efficacy in preclinical models fail in Phase 2/3 due to poor patient stratification, off-target toxicity, or lack of predictive biomarkers.

These failures are not merely scientific — they represent enormous economic waste and, more critically, delayed access to effective therapies for patients with unmet medical needs.
AI/ML offers a systematic solution to each of these pain points.

3. Objectives

O1. **Primary Objective 1:** Develop an ML pipeline (gradient boosting + graph neural networks) to identify and rank drug targets from integrated multi-omics datasets (RNA-seq, ATAC-seq, proteomics) with >80% precision@10 on benchmark disease datasets.

O2. **Primary Objective 2:** Build a protein-ligand binding affinity prediction module using transformer-based models trained on ChEMBL and PDB datasets, achieving Pearson r > 0.85 on held-out test sets.

O3. **Primary Objective 3:** Develop a patient stratification classifier for clinical trial enrollment using electronic health record (EHR) data and biomarker profiles, reducing predicted trial failure due to population heterogeneity by ≥30%.

O4. **Secondary Objective:** Create an open-source, modular web interface (BioMind Portal) that allows non-ML researchers to upload omics datasets and receive target prioritization reports without programming expertise.

O5. **Secondary Objective:** Publish ≥2 peer-reviewed manuscripts and present findings at one international biotechnology or bioinformatics conference.

4. Methodology

4.1 Phase I — Data Integration and Target Discovery (Months 1–6)

We will integrate publicly available multi-omics datasets from TCGA, GTEx, GEO, UniProt, and ENCODE using standardized bioinformatics pipelines (Snakemake workflows). Data harmonization will address batch effects using ComBat-seq and Harmony algorithms. A heterogeneous biological knowledge graph will be constructed using Neo4j, encoding

relationships between genes, proteins, diseases, pathways, and drugs derived from STRING, Reactome, and DrugBank.

Target identification will employ a Graph Attention Network (GAT) trained on known disease-gene associations from OMIM and DisGeNET. The model will propagate biological signals through the knowledge graph to score candidate targets by disease relevance, druggability (structural data from AlphaFold), and novelty. Output: a ranked list of top-50 candidate targets per disease query.

4.2 Phase I — Protein-Ligand Affinity Prediction

A transformer-based deep learning model (inspired by DeepDTA and MolBERT architectures) will be trained on >2 million protein-ligand pairs from ChEMBL 33 and BindingDB. Protein sequences will be encoded using ESM-2 embeddings; ligands will be represented as SMILES strings processed through a molecular graph transformer. The model will output predicted binding affinity (pKd/pKi values) and ADMET (absorption, distribution, metabolism, excretion, toxicity) properties.

4.3 Phase II — Clinical Trial Patient Stratification (Months 7–12)

Using de-identified EHR data (partnering with academic medical centers under IRB-approved protocols), we will train an XGBoost + Shapley Additive Explanations (SHAP) classifier to predict patient response to targeted therapies based on genomic biomarkers, comorbidities, and prior treatment history. Model interpretability will be prioritized via SHAP force plots to ensure clinical utility and regulatory compliance (FDA AI/ML guidance, 2021).

4.4 Validation Strategy

All models will be validated using 5-fold cross-validation and external test sets from independent cohorts not used in training. Target prioritization outputs will be benchmarked against gold-standard drug targets in approved oncology drugs. Wet lab partners will validate 3–5 top-ranked novel targets using CRISPR knockout experiments in relevant cell lines.
Clinical stratification models will be benchmarked against historical trial enrollment data.

5. Expected Outcomes

- ✓ A validated, open-source BioMind software framework deployable on cloud (AWS/GCP) oron-premise HPC clusters.
- ✓ Multi-omics target prioritization achieving >80% precision@10 on 3 disease benchmarks(cancer, neurodegeneration, metabolic disease).
- ✓ Protein-ligand affinity model with Pearson r > 0.85 on ChEMBL held-out test set, enablingvirtual screening of >10 million compound libraries.

✓ A patient stratification classifier reducing predicted misclassification in trial enrollment by≥30% compared to standard-of-care clinical criteria.

✓ An open-access BioMind Portal web application for non-programmer researchers,reducing time from data to biological insight from months to hours.

✓ Minimum 2 peer-reviewed publications in journals such as Nature Methods,Bioinformatics, or Cell Systems.

✓ A proof-of-concept demonstration validating ≥1 novel AI-predicted drug target in a cellular disease model.

Broader Impact: If BioMind reduces drug discovery preclinical timelines by even 20%, it could save an estimated $500M per drug program and bring life-saving therapies to patients 1–2 years earlier. The open-source model ensures equitable access for researchers in low- and middle-income countries who lack expensive proprietary platforms.

6. Project Timeline

Phase	Months	Key Activities	Deliverable
Phase I-A	1–2	Data curation, harmonization, knowledge graph construction	Integrated multi-omics database
Phase I-B	3–4	GAT model development, target prioritization training	Target ranking pipeline v1.0
Phase I-C	5–6	Protein-ligand model training, ADMET prediction module	Binding affinity predictor + preprint
Phase II-A	7–8	EHR data access, preprocessing, feature engineering	Patient stratification dataset
Phase II-B	9–10	Classifier training, SHAP explainability, validation	Stratification model + validation report
Phase II-C	11–12	BioMind Portal development, integration testing, publication	Web platform + peer-reviewed paper

7. Budget Overview

Budget Category	Estimated Cost (USD)	Justification

Personnel (2 FTE researchers, 1 year)	$60,000	Bioinformatician + ML Engineer
Cloud Computing (AWS/GCP GPU instances)	$15,000	Model training, large-scale inference
Database Licenses & Software	$5,000	ChEMBL, PDB API, Neo4j Enterprise
Wet Lab Validation (CRISPR, cell lines)	$12,000	Validation of top-3 predicted targets
Conference Presentations & Publication	$5,000	Open-access fees, travel
Miscellaneous / Contingency (10%)	$8,000	Unexpected costs
TOTAL	**~$105,000**	

8. Ethical Considerations and Risk Mitigation

All human EHR data used in the patient stratification module will be fully de-identified in compliance with HIPAA (US) and GDPR (EU) standards. IRB approval will be obtained from partner institutions prior to data access. The BioMind platform will implement role-based access controls and audit logs. Model outputs will be clearly labeled as decision-support tools, not diagnostic conclusions, to prevent inappropriate clinical reliance on AI predictions without physician oversight.

Algorithmic bias will be proactively assessed: training datasets will be audited for racial, ethnic, and sex representation, and separate performance metrics will be reported across demographic subgroups. Interpretability (via SHAP) will be mandatory for any clinical-facing model, ensuring that AI recommendations can be interrogated by clinicians and regulators.

All code, trained model weights, and processed datasets (where permitted by data use agreements) will be released under an MIT open-source license on GitHub and deposited in Zenodo to ensure reproducibility and broad scientific access.

9. Conclusion

BioMind represents a strategic investment in the infrastructure of modern drug discovery. By building an AI/ML platform purpose-designed for multi-omics biotechnology research, this project addresses three of the most costly and time-consuming bottlenecks in translational medicine. The modular, open-source design ensures that advances are democratized — available to researchers in universities and institutes in developing nations, not only well-funded commercial laboratories.

AI is not a replacement for biological intuition or experimental science. Rather, it is a tool that extends the reach of researchers — enabling them to ask bolder questions, test more hypotheses, and identify therapeutic opportunities that would remain invisible within datasets too complex for human analysis alone. BioMind is a step toward a future where the gap between biological data and clinical therapy narrows from decades to years.

References

Barabási, A. L., Gulbahce, N., & Loscalzo, J. (2011). Network medicine: A network-based approach to human disease. *Nature Reviews Genetics, 12*(1), 56–68.

Chen, H., Engkvist, O., Wang, Y., Olivecrona, M., & Blaschke, T. (2018). The rise of deep learning in drug discovery. *Drug Discovery Today, 23*(6), 1241–1250.

DiMasi, J. A., Grabowski, H. G., & Hansen, R. W. (2016). Innovation in the pharmaceutical industry: New estimates of R&D; costs. *Journal of Health Economics, 47*, 20–33.

FDA. (2021). Artificial intelligence and machine learning (AI/ML)-based software as a medical device (SaMD) action plan. U.S. Food and Drug Administration.

Jumper, J., Evans, R., Pritzel, A., et al. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature, 596*(7873), 583–589.

Lin, Z., Akin, H., Rao, R., et al. (2023). Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science, 379*(6637), 1123–1130. [ESM-2]

Lotfi Shahreza, M., Ghadiri, N., Mousavi, S. R., et al. (2018). A review of network-based approaches to drug repositioning. *Briefings in Bioinformatics, 19*(5), 878–892.

Nguyen, T., Le, H., Quinn, T. P., et al. (2021). GraphDTA: Predicting drug–target binding affinity with graph neural networks. *Bioinformatics, 37*(8), 1140–1147.

Stokes, J. M., Yang, K., Swanson, K., et al. (2020). A deep learning approach to antibiotic discovery. *Cell, 180*(4), 688–702.

Topol, E. J. (2019). High-performance medicine: The convergence of human and artificial intelligence. *Nature Medicine, 25*(1), 44–56.

Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. *Advances in Neural Information Processing Systems, 30.*